

LLM PITFALLS			SEVERITY <small>LOW ▼ ● ▲ HIGH</small>	0.0%	31.0%	58.6%	31.0%	58.6%	37.9%
CATEGORY	PITFALL		MITIGATION	GPT 3.5	GPT 4	GPT 4-MINI	CLAUDE 3.5	CLAUDE 4	SONAR
Design	Using generative AI as a first resort	▲	Consider more explainable, scrutable tech	Not evaluated	Not evaluated	Not evaluated	Not evaluated	Not evaluated	Not evaluated
	No baseline or state-of-the-art performance	●	Measure SOTA performance	Not evaluated	Not evaluated	Not evaluated	Not evaluated	Not evaluated	Not evaluated
	No success criterion	●	Ask user/customer what success means	Not evaluated	Not evaluated	Not evaluated	Not evaluated	Not evaluated	Not evaluated
Instruction following	Refusal to answer on unreasonable grounds	▼	Add context, remove ambiguity	Fail	Pass	Pass	Fail	Pass	Pass
	Misinterpretation, failure to clarify	▲	Instruct to clarify, add context, be explicit	Fail	Pass	Pass	Fail	Pass	Pass
	Failure to reject false or inconsistent query	●	Avoid incorrect or inconsistent prompts	Fail	Fail	Fail	Pass	Pass	Pass
	Arbitrary execution of instructions in data	▲	Explicitly separate instructions and data	Fail	Fail	Fail	Fail	Pass	Fail
	Difficulty manipulating sub-token elements	●	Pre-process sub-token elements	Fail	Fail	Pass	Fail	Fail	Fail
	Incorrectly formatted output	●	Provide format schema or template	Fail	Fail	Fail	Fail	Pass	Fail
Ambiguity & coherence	Ambiguity in the response	▲	Prompt engineering, pre-check responses	Fail	Fail	Pass	Fail	Fail	Fail
	Incoherent rambling or glitching	▼	Reword prompt and retry	Fail	Fail	Pass	Not evaluated	Pass	Fail
	Inconsistent reasoning across conversation	●	Shorten conversations, provide recaps	Fail	Pass	Pass	Not evaluated	Pass	Pass
	Irrelevance	▼	Prompt engineering, pre-check responses	Fail	Pass	Pass	Pass	Pass	Fail
Response quality	Simplistic (correct but poor) responses	●	Role allocation, few-shot prompts	Fail	Pass	Pass	Pass	Fail	Pass
	Incompleteness	●	Role allocation, few-shot prompts	Fail	Fail	Pass	Fail	Fail	Fail
	Faulty reasoning	▲	Chain-of-thought, tools, few-shot prompts	Fail	Pass	Pass	Pass	Pass	Fail
	Faulty premises	▲	Explicit chain-of-thought, break steps down	Fail	Fail	Fail	Fail	Fail	Fail
	Prompt sensitivity	●	Add more context and more instruction	Fail	Fail	Pass	Pass	Pass	Fail
	Overthinking	●	Provide more context about problem domain	Fail	Fail	Pass	Fail	Fail	Fail
	Overfit to similar task	▼	Emphasize differences from overfitted case	Fail	Fail	Fail	Fail	Pass	Fail
	Distraction	▼	Use neutral wording	Fail	Fail	Fail	Fail	Fail	Fail
	Verbose	▼	Specify minimal requirements	Fail	Fail	Fail	Fail	Fail	Fail
Factual errors	Incorrect factual recall	▲	RAG, knowledge graphs, pre-check response	Fail	Fail	Pass	Fail	Pass	Fail
	Fabricated or non-useful references	▲	RAG, knowledge graphs	Fail	Fail	Fail	Fail	Fail	Pass
	Coherent fiction	▲	RAG, knowledge graphs, pre-check response	Fail	Fail	Fail	Pass	Pass	Pass
	Out-of-date responses	●	RAG, tools	Fail	Pass	Fail	Fail	Fail	Pass
	Overconfidence in certainty	●	Disallow expressions of uncertainty	Fail	Fail	Pass	Fail	Fail	Fail
Bias & ethics	Explicit bias	▲	Improve alignment and system prompt	Fail	Fail	Pass	Pass	Pass	Pass
	Implicit bias	▲	Improve alignment, avoid leading signals	Fail	Fail	Fail	Fail	Pass	Fail
	Toxic or offensive responses	▲	Pre-check response	Fail	Fail	Fail	Fail	Pass	Fail
	Copyright infringement	▲	Knowledge graphs	Fail	Pass	Pass	Pass	Pass	Pass
	Cultural or ideological bias	●	Fine-tuning, alignment	Fail	Pass	Pass	Pass	Fail	Pass
Application	No transparency about role of AI	▲	Be clear	Not evaluated	Not evaluated	Not evaluated	Not evaluated	Not evaluated	Not evaluated
	Not collecting signals from users	▼	Collect signals	Not evaluated	Not evaluated	Not evaluated	Not evaluated	Not evaluated	Not evaluated
	Not providing references	●	Add references	Not evaluated	Not evaluated	Not evaluated	Not evaluated	Not evaluated	Not evaluated
Implementation	No rigorous evaluation protocol	▲	Create testing and reporting pipeline	Not evaluated	Not evaluated	Not evaluated	Not evaluated	Not evaluated	Not evaluated
	Low AI literacy among users	●	Education and culture	Not evaluated	Not evaluated	Not evaluated	Not evaluated	Not evaluated	Not evaluated
	Overreliance on model responses	▲	Education, avoiding risky applications	Not evaluated	Not evaluated	Not evaluated	Not evaluated	Not evaluated	Not evaluated
	Overreliance on human-in-the-loop	●	Consultation, training, improved interface	Not evaluated	Not evaluated	Not evaluated	Not evaluated	Not evaluated	Not evaluated
	No user training	●	Training	Not evaluated	Not evaluated	Not evaluated	Not evaluated	Not evaluated	Not evaluated
	No consideration of security or ethics	▲	Proper governance and oversight	Not evaluated	Not evaluated	Not evaluated	Not evaluated	Not evaluated	Not evaluated